

```
import pandas as pd
import numpy as np
```

```
df = pd.DataFrame([[8, 8, 1], [7, 9, 1], [6, 10, 0], [5, 12, 0]], columns = ['CGPA', 'ATS Score', 'Placed'])
df
```

	CGPA	ATS Score	Placed	
0	8	8	1	
1	7	9	1	
2	6	10	0	
3	5	12	0	

Next steps: [Generate code with df](#) [View recommended plots](#) [New interactive sheet](#)

```
def initialize_parameters(layer_dimensions):
    # Set a fixed random seed to make the results reproducible
    np.random.seed(3)

    # Create an empty dictionary to store the weight and bias parameters
    parameters = {}

    # Get the number of layers in the network (including input and output layers)
    L = len(layer_dimensions)

    # Loop through layers 1 to L-1 (we skip layer 0 because it's the input layer)
    for l in range(1, L):
        # Initialize weights for layer l with small values (here all 0.1 for simplicity)
        # Shape: (number of neurons in previous layer, number of neurons in current layer)
        parameters['W' + str(l)] = np.ones((layer_dimensions[l - 1], layer_dimensions[l])) * 0.1

        # Initialize biases for layer l with zeros
        # Shape: (number of neurons in current layer, 1)
        parameters['b' + str(l)] = np.zeros((layer_dimensions[l], 1))

    # Return the dictionary containing all the initialized parameters
    return parameters
```

```
# Utility Functions
def sigmoid(Z):

    A = 1/(1+np.exp(-Z))

    return A
```

```
def linear_forward(A_prev, W, b):
    # A_prev: activations from the previous layer (shape: size of previous layer x number of examples)
    # W: weights matrix for the current layer (shape: size of previous layer x size of current layer)
    # b: bias vector for the current layer (shape: size of current layer x 1)

    # Compute the linear transformation Z = W^T * A_prev + b
    # W.T gives shape: (size of current layer x number of examples)
    Z = np.dot(W.T, A_prev) + b

    # Apply sigmoid activation to the linear output Z
    A = sigmoid(Z)

    # Return the activation output A to be passed to the next layer
    return A
```

```
# Main Forward Propagation function for a single row instance (i.e., one data sample)
def L_layer_forward(X, parameters):
    # X: input vector of shape (input size, 1) - i.e., a single column (one data instance)
    # parameters: dictionary containing all the weights and biases for each layer

    A = X # Initial activation is the input feature vector
    L = len(parameters) // 2 # Number of layers in the network (each has W and b)

    # Loop through all the layers in the network
    for l in range(1, L + 1):
        A_prev = A # Store previous activation for computing current layer's activation

        # Retrieve current layer's weights and bias
        Wl = parameters['W' + str(l)]
        bl = parameters['b' + str(l)]

        # Perform linear forward computation: Z = W.T @ A_prev + b
        # NOTE: Your 'W' is stored as shape (prev_layer, current_layer), so we need W.T
        A = linear_forward(A_prev, Wl, bl)

    # Optional debugging prints:
    print("A" + str(l-1) + ": ", A_prev)
    print("W" + str(l) + ": ", Wl)
```

```

print("b" + str(l) + ": ", b1)
print("A" + str(l) + ": ", A)
print("--" * 20)

```

```

# Return final output (last activation), and the last activation before it
return A, A_prev

```

```

X = df[['CGPA', 'ATS Score']].values[0].reshape(2,1) # Shape(no of features, no. of training example)
y = df[['Placed']].values[0][0]

```

```

# Parameter initialization
parameters = initialize_parameters([2,2,1])

```

```

y_hat,A1 = L_layer_forward(X,parameters)
y_hat = y_hat[0][0]

```

```

update_parameters(parameters,y,y_hat,A1,X)

```

```

print('Loss for this student - ',-y*np.log(y_hat) - (1-y)*np.log(1-y_hat))

```

```

parameters

```

```

↺ A0: [[8]
      [8]]
W1: [[0.1 0.1]
     [0.1 0.1]]
b1: [[0.]
     [0.]]
A1: [[0.83201839]
     [0.83201839]]
-----
A1: [[0.83201839]
     [0.83201839]]
W2: [[0.1]
     [0.1]]
b2: [[0.]
     [0.]]
A2: [[0.54150519]]
-----
Loss for this student - 0.613402628898913
{'W1': array([[0.10000513, 0.10000513],
              [0.10000513, 0.10000513]]),
 'b1': array([[6.41054186e-07],
              [6.41054186e-07]]),
 'W2': array([[0.10003815],
              [0.10003815]]),
 'b2': array([[4.5849481e-05]])}

```

```

# Above:
# A0 = [8, 8] – This is the input vector (features from the first row of the dataset)
# W1 = Weights for Layer 1 – These connect the input layer to the first hidden layer
# b1 = [b11, b12] – Biases for the neurons in hidden layer 1

```

```

# These are passed to linear_forward(A_prev, W1, b1) which computes:
# Z1 = W1.T @ A0 + b1
# A1 = Z1 in this case (since there's no activation function yet), so:
# A1 = [011, 012] – outputs from the first hidden layer (linear combinations)

```

```

# Next layer:
# A1 is now passed as input to the next layer with W2, b2 to get:
# A2 = [021] – the output from the final layer (the prediction)

```

```

# The function L_layer_forward returns:
# - A2: the final output (predicted Package)
# - A1: the activation from the previous layer (needed for backprop)

```

```

# WHY RETURN A_prev TOO?
# During backpropagation, we compute gradients like:
#  $\partial L / \partial W = A_{prev}.T @ dZ$ 
# That's why we need the activations of the previous layer (A_prev or Oij values)
# They're essential for calculating the gradients during the backward pass.

```

```

def update_parameters(parameters, y, y_hat, A1, X):
    # -----
    # Output layer (layer 2) updates
    # -----

    # Gradient of loss w.r.t W2[0][0] = (dL/dy_hat) * (dy_hat/dz2) * (dz2/dW2[0]) = (y - y_hat) * A1[0]
    parameters['W2'][0][0] += 0.0001 * (y - y_hat) * A1[0][0]

    # Gradient of loss w.r.t W2[1][0] = (y - y_hat) * A1[1]
    parameters['W2'][1][0] += 0.0001 * (y - y_hat) * A1[1][0]

    # Bias update for output layer: gradient = (y - y_hat)
    parameters['b2'][0][0] += 0.0001 * (y - y_hat)

    # -----
    # Hidden layer (layer 1) updates
    # -----

    # For neuron 1 in hidden layer:

```

```
A0:  [[7]
[9]]
w1:  [[0.10000513 0.10000513]
[0.10000513 0.10000513]]
b1:  [[6.41054186e-07]
[6.41054186e-07]]
A1:  [[0.83202994]
[0.83202994]]
-----
A1:  [[0.83202994]
[0.83202994]]
w2:  [[0.10003815]
[0.10003815]]
b2:  [[4.5849481e-05]
[0.54153291]]
-----
Loss for this student - 0.6133514436691428
{'W1': array([[0.10000962, 0.1000109 ],
               [0.10000962, 0.1000109 ]]),
 'b1': array([[1.28227883e-06],
               [1.28227883e-06]]),
 'W2': array([[0.10007629],
               [0.10007629]]),
 'b2': array([[9.16961903e-05]]]}
```

```
A0: [[ [6]
[10]]
W1: [[ [0.10000962 0.1000109 ]
[0.10000962 0.1000109 ]]
b1: [[ [1.28227883e-06]
[1.28227883e-06]]
A1: [[ [0.83204007]
[0.83204294]]
-----
A1: [[ [0.83204007]
[0.83204294]]
W2: [[ [0.10007629]
[0.10007629]]
b2: [[ [9.16961903e-05]]
[9.16961903e-05]]
A2: [[ [0.54156062]]
[0.54156062]]
-----
Loss for this student - 0.7799272184937318
{'W1': array([[0.10000507, 0.10000333],
[0.10000507, 0.10000333]]),
'b1': array([[5.25214767e-07],
[5.25225084e-07]]),
'W2': array([[0.10003123],
[0.10003123]])}
```

```
    [0.10003123]]],
    'b2': array([[3.75401279e-05]])}
```

```
update_parameters(parameters,y,y_hat,A1,X)
parameters
```

```
↩ { 'W1': array([[0.10000053, 0.09999576],
               [0.10000053, 0.09999576]]),
    'b1': array([[-2.31508272e-07],
               [-2.31487641e-07]]),
    'W2': array([[0.09998617],
               [0.09998617]]),
    'b2': array([[-1.66159345e-05]])}
```

```
# Extract input features for the 4th training example (index 2)
X = df[['CGPA', 'ATS Score']].values[3].reshape(2,1) # Shape(no of features, no. of training example)
y = df[['Placed']].values[3][0]
```

```
y_hat,A1 = L_layer_forward(X,parameters)
y_hat = y_hat[0][0]
```

```
update_parameters(parameters,y,y_hat,A1,X)
```

```
print('Loss for this student - ',-y*np.log(y_hat) - (1-y)*np.log(1-y_hat))
```

```
parameters
```

```
↩ A0: [[ 5]
      [12]]
W1: [[0.10000053 0.09999576]
     [0.10000053 0.09999576]]
b1: [[-2.31508272e-07]
     [-2.31487641e-07]]
A1: [[0.84553589]
     [0.84552529]]
-----
A1: [[0.84553589]
     [0.84552529]]
W2: [[0.09998617]
     [0.09998617]]
b2: [[-1.66159345e-05]
     [0.54216614]]
-----
Loss for this student - 0.7812489129203491
{ 'W1': array([[0.099997 , 0.09998727],
               [0.099997 , 0.09998727]]),
  'b1': array([[-9.39181566e-07],
               [-9.39200617e-07]]),
  'W2': array([[0.09994033],
               [0.09994033]]),
  'b2': array([[-7.08325485e-05]])}
```

```
# Till now, the whole dataset got trained for 1 time(1 epoch)
# Let's start with epoch implementation
```

```
# Outer loop: iterate through entire dataset multiple times (epochs)
```

```
parameters = initialize_parameters([2,2,1])
epochs = 50
```

```
for i in range(epochs):
```

```
    Loss = []
```

```
    for j in range(df.shape[0]):
```

```
        X = df[['CGPA', 'ATS Score']].values[j].reshape(2,1) # Shape(no of features, no. of training example)
        y = df[['Placed']].values[j][0]
```

```
        # Parameter initialization
```

```
        y_hat,A1 = L_layer_forward(X,parameters)
        y_hat = y_hat[0][0]
```

```
        update_parameters(parameters,y,y_hat,A1,X)
```

```
        Loss.append(-y*np.log(y_hat) - (1-y)*np.log(1-y_hat))
```

```
    print('Epoch - ',i+1,'Loss - ',np.array(Loss).mean())
```

```
parameters
```



```

-----
A1: [[0.83219719]
      [0.83146478]]
W2: [[0.09932622]
      [0.09932647]]
b2: [[-0.00076653]]
A2: [[0.54102728]]
-----
A0: [[ 6]
      [10]]
W1: [[0.1000849  0.09975928]
      [0.10008503 0.09975881]]
b1: [[-7.57986149e-06]
      [-7.59695223e-06]]
A1: [[0.83220728]
      [0.83147768]]
-----
A1: [[0.83220728]
      [0.83147768]]
W2: [[0.09936441]
      [0.09936463]]
b2: [[-0.00072064]]
A2: [[0.54105502]]
-----
A0: [[ 5]
      [12]]
W1: [[0.1000804  0.09975178]
      [0.10008051 0.09975128]]
b1: [[-8.33023947e-06]
      [-8.34993427e-06]]
A1: [[0.84571225]
      [0.84498093]]
-----
A1: [[0.84571225]
      [0.84498093]]
W2: [[0.09931939]
      [0.09931965]]
b2: [[-0.00077474]]
A2: [[0.54168901]]
-----
Epoch - 50 Loss - 0.6969136669930831
{'W1': array([[0.10007689, 0.09974336],
              [0.10007699, 0.09974283]]),
 'b1': array([[ -9.03191730e-06,
              [-9.05433049e-06]]),
 'W2': array([[0.09927357],
              [0.09927387]]),
 'b2': array([[ -0.00082891]])}

```

```

# After all epochs, print the final learned parameters
# These include W1, b1, W2, b2 after being refined over 50 epochs
parameters

```

```

➦ {'W1': array([[0.10007689, 0.09974336],
                [0.10007699, 0.09974283]]),
    'b1': array([[ -9.03191730e-06,
                  [-9.05433049e-06]]),
    'W2': array([[0.09927357],
                  [0.09927387]]),
    'b2': array([[ -0.00082891]])}

```

Start coding or [generate](#) with AI.